


# Lecture 09: SVM, kNN, XGBoost, and Materials Discovery



Instructor: Sergei V. Kalinin

# Tasks, workflows, and workflow planning




**Task** can be classification, regression, clustering, optimization.  
**Algorithm** can be kNN, perceptron, DCNN, etc.


We can use complex algorithms for “simple” tasks and simple algorithms for complex tasks.



# Classification problem



# Support Vector Machines





**Perceptron:** minimize misclassification errors.

**SVM:** maximize the margin.


The margin is the distance between the separating hyperplane (decision boundary) and the training examples that are closest to this hyperplane, which are called **support vectors**.

# Regularization in SVMs



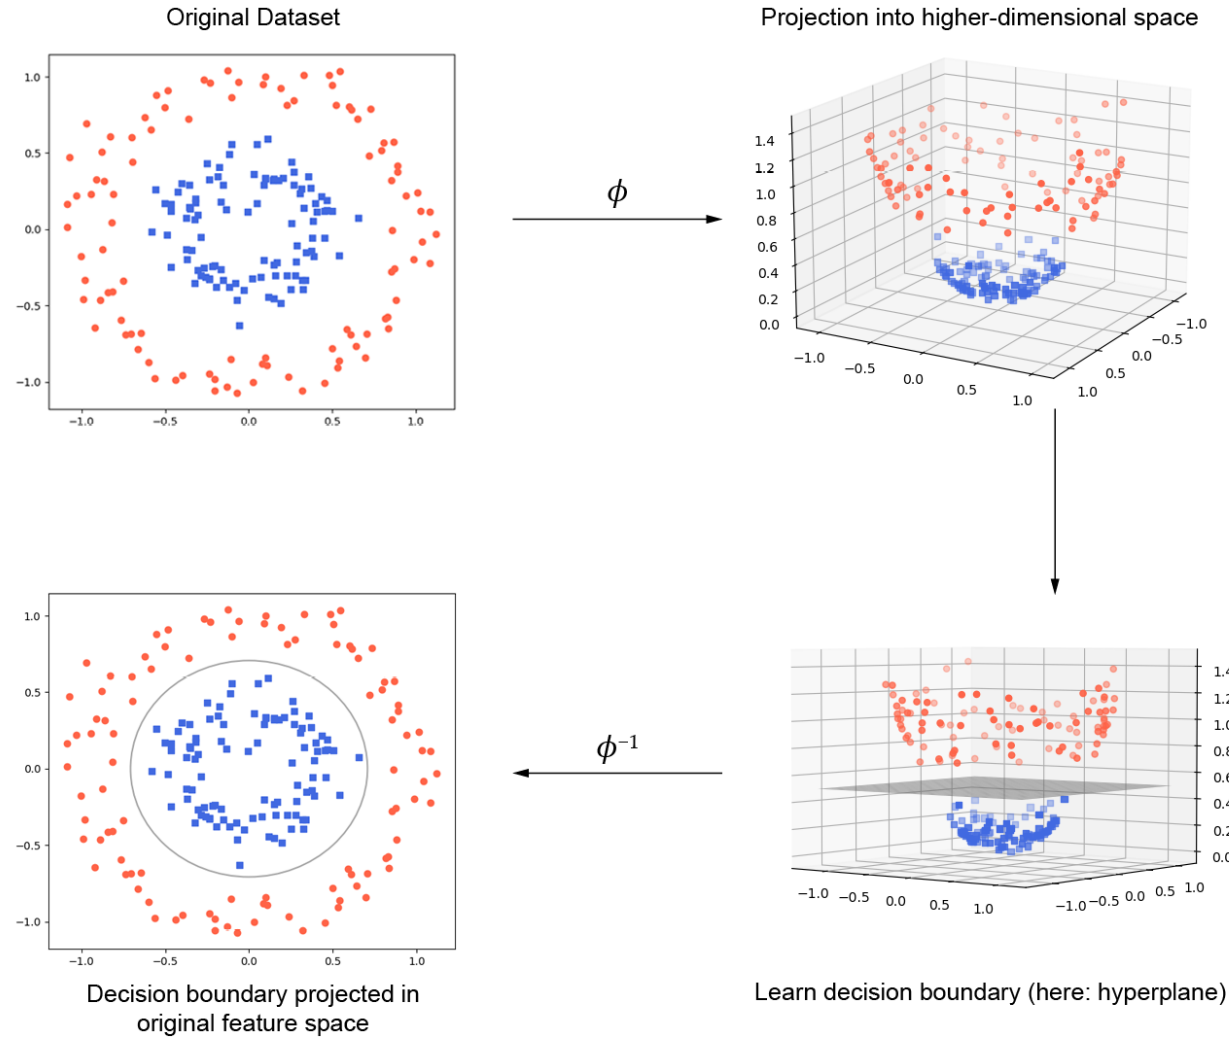
Large values of  $C$  correspond to large error penalties, whereas we are less strict about misclassification errors if we choose smaller values for  $C$ .

# Kernel SVM: Motivation



Non-linearly separable problem

# Kernel SVM



The idea behind **kernel methods** for dealing with linearly inseparable data is to create nonlinear combinations of the original features to project them onto a higher-dimensional space via a mapping function,  $\phi$ , where the data becomes linearly separable.

$$\phi(x_1, x_2) = (z_1, z_2, z_3) = (x_1, x_2, x_1^2 + x_2^2)$$

# Kernel SVM

- To train an SVM, in practice, we need to replace the dot product  $\mathbf{x}^{(i)T}\mathbf{x}^{(j)}$  by  $\phi(\mathbf{x}^{(i)})^T \phi(\mathbf{x}^{(j)})$  .
- To save the expensive step of calculating this dot product between two points explicitly, we define a **kernel function**  $\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \phi(\mathbf{x}^{(i)})^T \phi(\mathbf{x}^{(j)})$
- 
- One of the most widely used kernels is the **radial basis function (RBF)** kernel, or the **Gaussian kernel**:


$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \exp\left(-\gamma \|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|^2\right)$$

$\gamma$  is a free parameter to be optimized.

- Roughly speaking, the term “kernel” can be interpreted as a **similarity function** between a pair of examples.



# k Nearest Neighbours (kNN) Classifier



1. Choose the number of  $k$  and a distance metric
2. Find the  $k$ -nearest neighbors of the data record that we want to classify
3. Assign the class label by majority vote

KNN is an example of a **lazy learner**. It is called “lazy” not because of its apparent simplicity, but because it doesn’t learn a discriminative function from the training data but memorizes the training dataset instead.

# Parametric vs. Non-parametric methods

**Parametric models:** we estimate parameters from the training dataset to learn a function that can classify new data points without requiring the original training dataset anymore. Typical examples of parametric models are the perceptron, logistic regression, and the linear SVM.

**Non-parametric models** can't be characterized by a fixed set of parameters, and the number of parameters changes with the amount of training data. Two examples of non-parametric models are the decision tree classifier/random forest and the kernel (but not linear) SVM.

KNN belongs to a subcategory of non-parametric models described as **instance-based learning**. Models based on instance-based learning are characterized by memorizing the training dataset, and lazy learning is a special case of instance-based learning that is associated with no (zero) cost during the learning process.

From S. Raschka, Machine Learning with PyTorch and Scikit-Learn

# Pros and Cons of Memory Based Approaches

- The main advantage of memory-based approach is that the classifier immediately adapts as we collect new training data
- The downside is that the computational complexity for classifying new examples grows linearly with the number of examples in the training dataset in the worst-case scenario—unless the dataset has very few dimensions (features) and the algorithm has been implemented using efficient data structures for querying the training data more effectively.
- Such data structures include k-d tree ([https://en.wikipedia.org/wiki/K-d\\_tree](https://en.wikipedia.org/wiki/K-d_tree)) and ball tree ([https://en.wikipedia.org/wiki/Ball\\_tree](https://en.wikipedia.org/wiki/Ball_tree)), which are both supported in scikit-learn. Furthermore, next to computational costs for querying data, large datasets can also be problematic in terms of limited storage capacities.
- However, in many cases when we are working with relatively small to medium-sized datasets, memory-based methods can provide good predictive and computational performance and are thus a good choice for approaching many real-world problems.

# Distances and Neighbours

- We need to choose a distance metric that is appropriate for the features in the dataset.
- Usual choice: Minkowski distance, a generalization of the Euclidean ( $p = 2$ ) and Manhattan ( $p = 1$ ) distance:

$$d(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \sqrt[p]{\sum_k |x_k^{(i)} - x_k^{(j)}|^p}$$

- Many other distance metrics are available in scikit-learn and can be provided to the **metric** parameter. They are listed at <https://scikit-learn.org/stable/modules/generated/sklearn.metrics.DistanceMetric.html>.
- KNN is very susceptible to overfitting due to the **curse of dimensionality**, where feature space becomes increasingly sparse for an increasing number of dimensions of a fixed-size training dataset.
- For regression and SVM, we use regularization to avoid this problem.
- We cannot use regularization for decision trees and KNN. Instead, we can use feature selection and dimensionality reduction techniques

# How materials are discovered?

Corning Ware glass was accidentally discovered via a furnace mishap



“The temperature gauge was stuck on 900 degrees,  
and I thought I had ruined the furnace ...

I grabbed some tongs to get it out as fast as I could,  
but the glass slipped out of the tongs and fell to the floor.

**The thing bounced and didn’t break.”**

Donald Stookey (1915-2014)

<https://cen.acs.org/articles/92/web/2014/12/Donald-StookeyGuy-Gave-Us-Corning.html>

<https://www.nytimes.com/2014/11/07/business/s-donald-stookey-inventor-of-corningware-dies-at-99.html>

Slide by Sterling Baird

# How materials are discovered?




Velcro photo, shape memory alloy licensed under CC BY | Sweet N Low, Teflon licensed under CC BY-SA  
“How are materials discovered?” by Taylor Sparks: <https://www.youtube.com/watch?v=RRNcqJSJ6vc>

Slide by Sterling Baird





# How materials are discovered?



<https://en.wikipedia.org/wiki/LK-99>

- 1986 –  $\text{YBa}_2\text{Cu}_3\text{O}_7$ . Gave rise to multiple families of Cu and Hg superconductors
- 2001 –  $\text{MgB}_2$ . Point compound
- 2006 – Layered iron pnictides. Gave rise to multiple families of superconductors

# How many molecules are there?



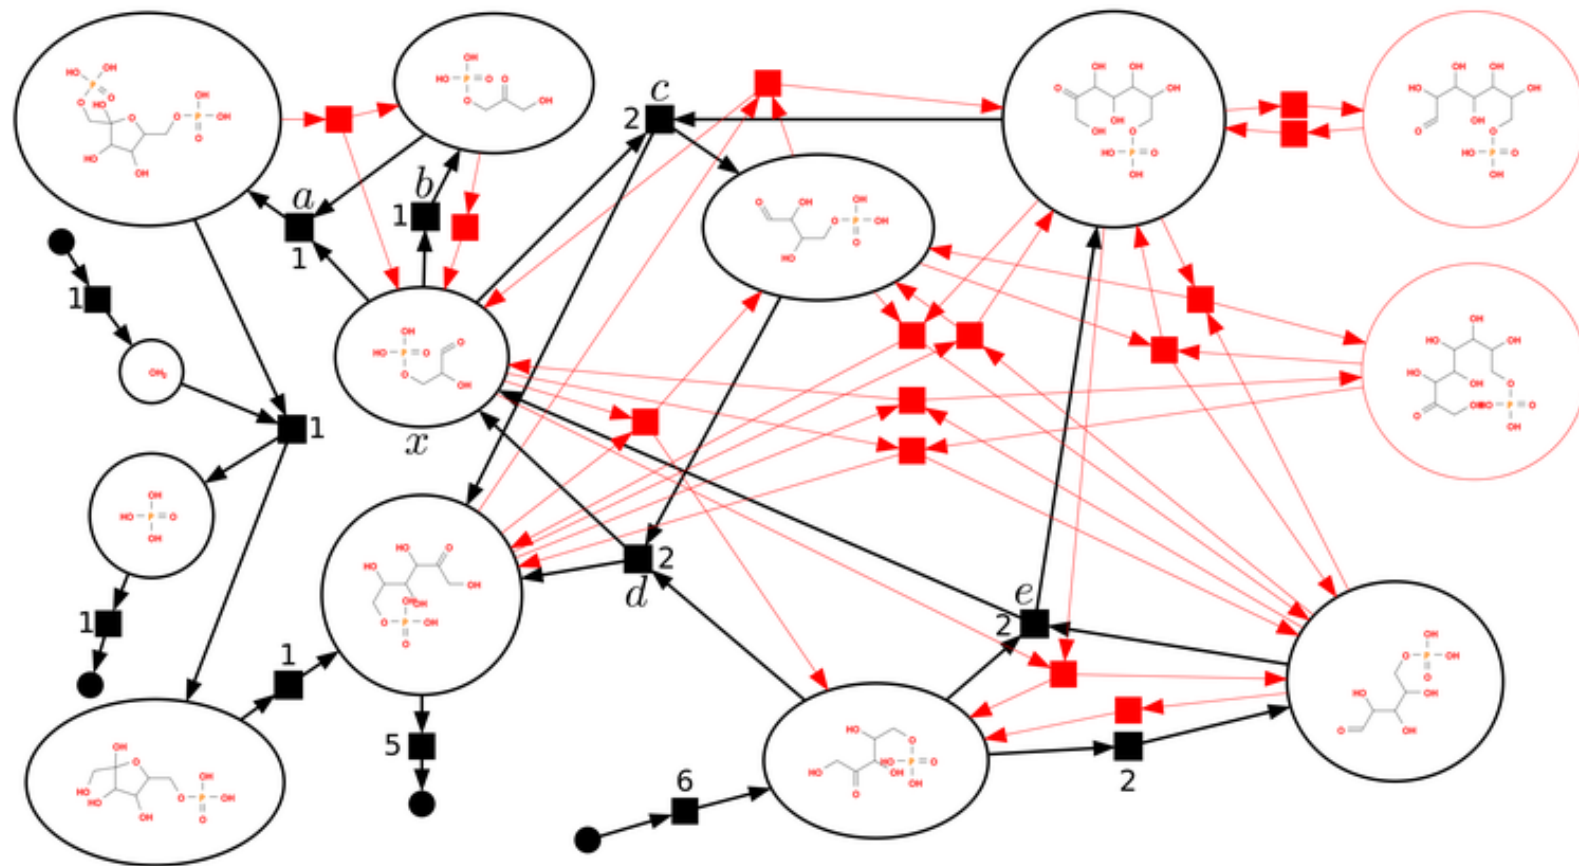
A chemical space often referred to in cheminformatics is that of potential biologically active molecules. Its size is estimated to be in the order of  $10^{60}$  molecules. The estimate restricts the chemical elements used to be C, H, O, N and S. It further makes the assumption of a maximum of 30 atoms to stay below 500 Daltons, allows for branching and a maximum of 4 rings and arrives at an estimate of  $10^{63}$ .

<https://www.cell.com/trends/chemistry/fulltext/S2589-5974%2820%2930288-4>

[https://en.wikipedia.org/wiki/Chemical\\_space](https://en.wikipedia.org/wiki/Chemical_space)




# Chemical reactions networks:



- Molecular property predictions: are they **likely** to be useful?
- Synthesizability scores: what would it **probably** take to make them
- Reaction network mining and retrosynthesis: can we identify **possible** synthetic pathways?
- Optimization of specific reaction conditions and pathways: myopic and non-myopic

# Why dimensionality is a problem?



Let's think about it as a search problem:



- **Alloying:** need maintain composition  $\sim 1\%$
- **Doping:** need maintain composition  $\sim 10^{-6}$
- Grid search is out for  $D > 3$  (experiment)

# Why dimensionality is a problem?


- Suppose that we have data for 1000 students' performance (discretized scores of 0; 25; 50; 75; 100)% in 2 courses c1 and c2. Then in total there are  $5 \times 5 = 25$  different grade combinations.
- If the 1000 students are randomly distributed among each grade combination, then on average there are 40 students with each possible grade combination, which is a good enough sample to draw conclusions such as if, for a student, `grade(c1) 50` and `grade(c2) 75`, then that student is likely to be a Math major.
- Now suppose there are 4 courses, then the number of possible grades combination is  $5^4 = 625$ , and an average number of students per combination is 1:6. For 10 courses, this number reduces to 0:0001024. This means that almost all possible combinations are never observed.

# Why dimensionality is a problem?

- Suppose  $n$  points in  $X$  are chosen uniformly at random from  $[0; 1]^m$  ( $m$ -cube). For the query point  $q$  grow a hypercube around  $q$  to contain  $f$  fraction of points ( $k = f n$ ) in  $X$ . This cube (the search space for  $q$ ) grows very large (covering almost the whole input space) in large dimension.
- The expected length of the edge of the search cube  $E_m(f) = f^{1/m}$ , i.e. in 10d to get 10% points around  $q$  need cube with edge length 0.8 (which is 80% of the whole cube, the input space). Similarly, to get only 1% points one needs to extend the search cube by 0.63 units along each dimension





# Why dimensionality is a problem?



## Suppose we have 5000 points:

- In 1d we have to explore 0.001 on average to capture 5 NN
- In 2d, on average we must explore 0.031 units along both dimensions to get 5 nearest neighbors points (about 3% of the whole cube).
- In 3d, on average we must go 10% of the total (unit) length in each of the 3 dimensions
- In 4d, we must explore 17.7% of unit length
- In 10d, we must go 50.1% of unit length along each dimension

# Why dimensionality is a problem?



dim $m$	volume of $m$ -ball	volume of $m$ -cube	ratio
2	$\pi$	$2^2$	$\sim 0.785$
3	$\frac{4}{3}\pi$	$2^3$	$\sim 0.523$
4	$\frac{\pi^2}{2}$	$2^4$	$\sim 0.308$
6	$\frac{\pi^3}{6}$	$2^6$	$\sim 0.080$
$m$	$\frac{\pi^{m/2}}{m/2!}$	$2^m$	$\rightarrow 0$

# Why dimensionality is a problem?

However if a dataset exhibit this phenomenon that the issue has be overcome by getting a larger training set (exponential in  $m$ ). One way to look at this is as follows.

To cover  $[-1, 1]^m$  with  $B_{m,1}$ 's, the number of balls  $n$  must be

$$n \geq \frac{2^m}{V_m(1)} = \frac{2^m}{\pi^{m/2}/m/2!} = \frac{m/2! 2^m}{\pi^{m/2}} \stackrel{m \rightarrow \infty}{\sim} \sqrt{m\pi} \left( \frac{m2^{m/2}}{2\pi e} \right)^{m/2}$$

For  $m = 16$  (a very small number) this  $n$  is substantially larger than  $2^{58}$

- In higher dimensions all the volume is in 'corners'
- Points in high dimensional spaces are isolated (empty surrounding)
- The probability that a randomly generated point is within  $r$  radius of  $q$  approaches 0 as dimensionality increases
- The probability of a close nearest neighbor in a data set is very small

